

Finding Cut On Spectral-Hard Graphs Using Local Algorithms

Lichen Zhang

Computer Science Department
Carnegie Mellon University
Pittsburgh, PA, 15213
lichenz@andrew.cmu.edu

Gary L. Miller

Computer Science Department
Carnegie Mellon University
Pittsburgh, PA, 15213
glmiller@cs.cmu.edu

Andy Yang

Computer Science Department
Carnegie Mellon University
Pittsburgh, PA, 15213
andrewy1@andrew.cmu.edu

Tian Luo

Computer Science Department
Carnegie Mellon University
Pittsburgh, PA, 15213
tianl1@andrew.cmu.edu

May 11, 2020

Abstract

We study the graph cut problem under local settings: algorithms have the constraint that its run time must be proportional to the set it outputs, instead of the entire graph. Quality of cut is measured under conductance metric. We experiment modified versions of `Nibble` algorithm in [1] and local pagerank algorithm in [2] on graphs that are designed to be “hard” on these algorithms.

Keywords: graph cut, local algorithms

1 Introduction

With the fast growth in number of social media users and size of community network graph, it becomes more and more important to answer the following type of query efficiently: given a node in a community network graph, find a cluster around given node. Here, we measure the quality of cluster by conductance metric, intuitively, it measures the ratio of out-of-cluster connectivity and internal connectivity of the cluster. We also restrict our attention to a class of algorithms we refer to *local algorithms*: run time of these algorithms must be proportional to the size of cluster it outputs instead of the size of entire graph. We remark local algorithms are crucial to today’s applications, especially on large social network graphs such as Facebook Graph [3] and Pinterest Graph [4], these graphs typically have billions of nodes and clustering query will be made for every node in a daily basis [5], so any algorithm whose run time is proportional to size of entire graph, even linear, is prohibitive to be deployed. In contrast, these graphs generally have small and relatively high quality clusters, which means local algorithms will be more than feasible in solving this kind of problems.

Local graph clustering algorithms have a relatively short but successful history. Starting from the groundbreaking work by Spielman and Teng [1], various algorithms such as local pagerank [2] and evolving set based algorithm [6] have been studied. All these algorithms exploit Lovasz-Simonovits Theorem [7], which connects conductance of a set and evolution of random walk, and inherits a Cheeger-type bound [8] in their output cluster. This essentially means that the cluster they output might have a $O(\sqrt{\lambda_2})$ from the best-possible cluster, where λ_2 is the second smallest eigenvalue of normalized Laplacian matrix of the graph. Recently, Wang et al. [9] have proposed a novel algorithm that is based on preflow-relabel algorithm due to Goldberg and Tarjan [10], and they claim that their algorithm won't suffer from Cheeger's bound. However, they don't give a concrete proof that they can avoid such problem. Instead, unlike previous random walk based work, their analysis does not depend on Lovasz-Simonovits Theorem. We remark that, though all of above algorithms have great theoretical interests and practical usage, and local pagerank algorithm has been deployed in various applications, there's almost no empirical evidence in showing that random walk based algorithms suffer from Cheeger's bound, and Wang et al.'s algorithm does not have this problem.

We are interested in the following two problems: 1). Does random walk based algorithms really have Cheeger's bound embedded in them, and can we verify this through experiments? 2). Can we confirm that Wang et al.'s algorithm really solve the problem of Cheeger's bound least through experimental evidence? Especially, we setup the experiment on graphs which are supposed to have a poor performance on algorithms rely on Cheeger's bound and verify modified versions of Nibble [1], local pagerank [2] and Wang et al.'s flow algorithm [9] on it. This report will focus on experiments of random walk based algorithms, and present some novel algorithms on simulating random walk on these graphs. For flow based algorithm by Wang et al., we refer readers to report by Andy Yang. For more analysis on graph construction and performance evaluation of random walk based algorithms, we refer readers to report by Tian Luo.

2 Problem Setup & Notations

2.1 Conductance

Let $G = (V, E)$ be an undirected graph, let $d(v)$ denote the degree of $v \in V$. Suppose $S \subseteq V$, define $\text{vol}(S) := \sum_{v \in S} d(v)$, and $E(S, T) = \{\{u, v\} \in E : u \in S, v \in T\}$, the *conductance* of set S is defined as

$$\Phi(S) := \frac{|E(S, \bar{S})|}{\min\{\text{vol}(S), \text{vol}(\bar{S})\}}$$

In case $S = G$, we define $\Phi(G) = 1$. And the *conductance* of G is defined as

$$\Phi_G := \min_{S \subseteq V} \Phi(S)$$

The *local graph clustering problem* is defined as follows: given $v \in V$, find a subset $C \subseteq V$ such that $v \in C$, and

$$\Phi(C) = \min_{S \subseteq V, v \in S} \Phi(S)$$

2.2 Graph Laplacian and Cheeger's Inequality

The *graph Laplacian matrix*, L is defined as $L := D - A$, where D is the *degree matrix* of G , which is a diagonal matrix with each entry corresponds to degree of each vertex, and A is the *adjacency matrix* of G . The *normalized Laplacian matrix*, \tilde{L} , is defined as $\tilde{L} = D^{-1/2} L D^{-1/2}$, let λ_2 denote the

second smallest eigenvalue of \tilde{L} . We remark that $0 < \lambda_2 \leq 1$. The *Cheeger's inequality* [8] connects Φ_G to λ_2 , specifically, it says

$$\frac{\lambda_2}{2} \leq \Phi_G \leq \sqrt{2\lambda_2}$$

Notice lower bound and upper bound differ by a $O(\sqrt{\lambda_2})$ factor, in Cheeger's proof of the upper bound [8], he presented an algorithm to find a cut of conductance $\sqrt{2\lambda_2}$, which means there exists a class of algorithms that finds cut whose conductance is $O(\sqrt{\lambda_2})$ away from best cut. Theoretically, Nibble algorithm [1], local pagerank [2] and evolving set [6] all have such problem.

2.3 Cartesian Product of Graphs

Let $G = (V_G, E_G)$, $H = (V_H, E_H)$, we define the *Cartesian product* of G and H , denoted by $G \square H$, as follows: let A_G, A_H be adjacency matrices of G and H , then

$$A_{G \square H} := A_G \oplus A_H = A_G \otimes I_H + I_G \otimes A_H$$

Here \oplus denotes the Kronecker sum, \otimes denotes the Kronecker product and I_H, I_G denote the identity matrices of size $|V_H|, |V_G|$. We remark several properties of $L_{G \square H}$ and L_G, L_H . Let $\lambda(L)$ denote set

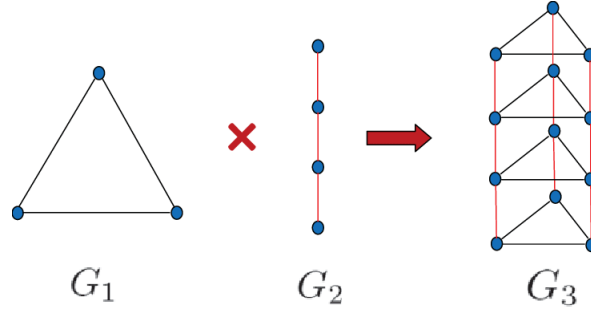


Figure 1: Example of Cartesian product of two graphs by [11]

of eigenvalues of L , then

$$\lambda(L_{G \square H}) = \{\lambda_i + \mu_j : \lambda_i \in \lambda(L_G), \mu_j \in \lambda(L_H)\}$$

In words, eigenvalues of $L_{G \square H}$ are direct sums of eigenvalues of L_G and L_H , and eigenvectors have similar structure: suppose x_i is an eigenvector of L_G and y_j is an eigenvector of L_H , then one of eigenvectors of $L_{G \square H}$ is computed as follows: 1). compute $x_i y_j^T$, i.e., outer product of two vectors, 2). flatten the matrix into a single vector. In words, eigenvectors of $L_{G \square H}$ can be viewed as direct products of eigenvectors of L_G, L_H . We further remind reader that since Laplacian matrix has rank $n - 1$, the smallest eigenvalue is 0, which means

$$\lambda_2(L_{G \square H}) = \min \{\lambda_2(L_G), \lambda_2(L_H)\}$$

2.4 Problem Setup

For the sake of simplicity, let $N \in \mathbb{N}_+$ such that $N = n^3$ for some $n \in \mathbb{N}_+$, consider a balanced binary tree T on $O(N^{2/3})$ nodes and a path P on $O(N^{1/3})$ nodes, it's worth noting that $\lambda_2(L_T) = \Theta(\frac{1}{N^{2/3}})$ and $\lambda_2(L_P) = \Theta(\frac{1}{N^{2/3}})$, so they are at the same order and differ by only a constant. Specifically, $\lambda_2(P) \leq \frac{12}{N^{2/3}}$ and $\lambda_2(T) \geq \frac{0.5}{N^{2/3}}$, so we can set $|V_P| = \lceil 4\sqrt{3} \rceil N^{1/3}$ and $|V_T| = N^{2/3}$, so that

$\lambda_2(L_T \square P) = \lambda_2(L_P)$. Intuitively, one can think of the graph as balanced tree lie from top to bottom, and each two consecutive trees are connected with path. Such eigenstructure will cause Cheeger-type to find a cut that cuts all paths between the two trees at center, which cuts $N^{2/3}$ edges, however, a much better cut will be cutting one edge near root for each tree, which only cuts $O(N^{1/3})$ edges. We expect algorithms like `Nibble` and local pagerank to fail on this example — i.e., they will find a cluster that cuts the graph horizontally, not vertically. On the other hand, we expect algorithm by Wang et al. [9] to perform well on this example. We refer to this particular $T \square P$ graph as *tree-cross-line graph*.

3 Main Results

We present two experiments using `Nibble` and local pagerank on tree-cross-line graph. Specifically, we implement modified versions of these algorithms, since our focus is to evaluate the quality of cut they output instead of getting a fast algorithm. We perform both experiment on a tree-cross-line graph with $N = 15625$, with tree size 625 and path size 25, notice the path size is $\lceil 4\sqrt{3} \rceil$ times smaller than the size that will make $\lambda_2(L_P)$ smaller (this is due to computation cost for matrix vector product or linear system solve for $N \times N$ matrices), so we expect these algorithms can find a decently good cut on our graph. However, the experiments show that the cut they produce is more of a horizontal cut, and this gives evidence that these random walk based algorithms are subject to Cheeger’s bound.

3.1 Experiment Setup

We implement modified versions of both `Nibble` and local pagerank.

3.1.1 Nibble as Random Walk

The vanilla `Nibble` algorithm presented by Spielman and Teng [1] is as follows:

```

INPUT: Graph G, starting node u, epsilon
let p be an indicator vector on starting node u;
let W be the random walk transition matrix;
for i = 1, 2, ..., max iterations do:
    p = Wp;    // perform walk
    let u be the vector such that:
        u[j] = p[j] / d(j), where d(j) is the degree of node j;
        set u[j] = 0 if u[j] < epsilon;
    sort u in descending order;
    examine all level cuts formed by values in sorted u,
    record the best one;
output the best level cut;

```

Notice the step that removes probability on entry where $\frac{p[j]}{d(j)}$ is too small only affects the run time of algorithm but not the quality of cut, in fact, without removing those small entries, the cut it outputs will have a better quality guarantee. Our modified version won’t remove small probability entries in u .

3.1.2 Local PageRank as Solving Laplacian System

The vanilla version of local pagerank algorithm presented by Anderson, Chung and Lang [2] is even simpler than Nibble:

```

INPUT: Graph G, starting node u, reset probability a
let p be a vector that is supposed to be the solution to a linear system;
let W be random walk transition matrix;
let m be an indicator vector on starting node u;
solve system  $[I - (1 - a)W]p = am$ ;
form vector u where  $u[j] = p[j] / d(j)$ ;
sort u in descending order;
examine all level cuts formed by values in sorted u;
output the best level cut;

```

We will then show can be done by one Laplacian solve.

Proposition 3.1. *The vector u can be computed directly via a Laplacian solve, specifically, u is the solution to the system*

$$(\beta D + L)u = \beta m$$

Where we assume $W = A^T D^{-1}$ to be the standard random walk transition matrix, and $\beta = \frac{a}{1-a}$.

We defer the proof to appendix section.

3.1.3 Experiment Environment & Parameter Setup

We conduct all experiments under Python 3.7, with numerical computation package `numpy` and graph package `networkx`.

Following our proceeding discussion, we implement `Nibble` by performing matrix vector product of W and p , this is achieved via `numpy`'s matrix multiplication, i.e., `@` operation. For local pagerank, we solve the Laplacian system

$$(\beta D + L)u = \beta m$$

Directly using `np.linalg.solve`.

For parameter selection, we choose $N = 15625$, which is 25^3 . We choose tree size to be $625 = 25^2$, and path size to be 25, it is worth noting that if we want to ensure $\lambda_2(P) < \lambda_2(T)$, we need path size to be 7 times larger. However, due to space constraint on the size of the matrix, we have to relax the condition by making the path smaller. This will make $\lambda_2(T) < \lambda_2(P)$, but the difference in magnitude is less than 10^{-3} .

3.2 Experiment Result & Analysis

In this section, we also include surprises, since the general results are quite surprise and a little bit of out of expectation. We first give the benchmark of our experiment:

	number of cut edges	conductance
vertical cut	25	0.0008
horizontal cut	625	0.0204

Table 1: Benchmarks for vertical and horizontal cuts

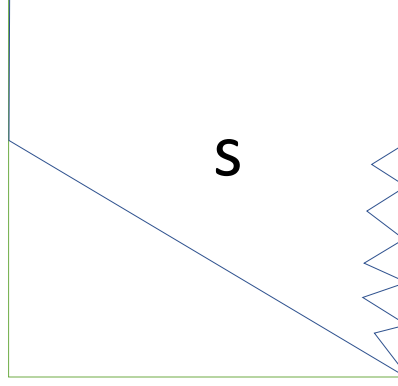


Figure 2: Shape of cut produced by Nibble

We remark that a balanced binary tree of size 625 has height 10. For both **Nibble** and local pagerank, we start from one of the leaves at the top tree.

For **Nibble**, we simulate random walk for 100 iterations, the resulting cut has size **961** edges and conductance **0.03140**. The shape of the cut is drawn as follows: Here S refers the cut that **Nibble** outputs, we draw it in a square fashion, where the vertical sides can be viewed as paths and horizontal sides can be viewed as trees. Intuitively, one can view it as we lie trees horizontally, from top to bottom on the square, so the horizontal cut should correspond to a horizontal line on the square, and a vertical cut should correspond to a vertical line on the square. As we can see from the figure, the cut is more of a horizontal cut, but it also cuts many edges on trees much lower on the path, so the cut looks like a diagonal cross. It also contains fewer and fewer nodes as it goes down the path, which is a sign that the cut it produces resembling the horizontal Cheeger's cut. More specifically, it contains the first 9 trees entirely, and most part of next 3 trees, and subsequent lower trees have fewer and fewer parts being in the cut. To better illustrate the idea, we draw the first 4 levels near root of the 17th tree:

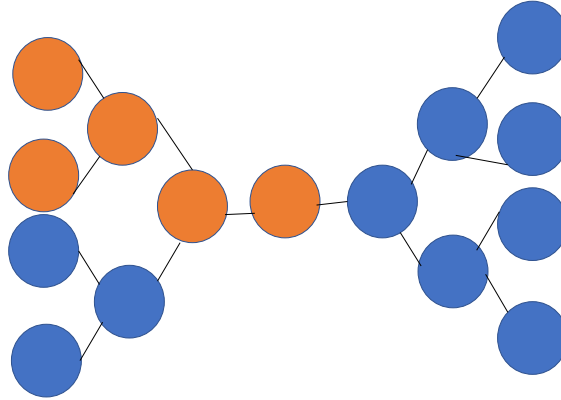


Figure 3: First 4 levels of 17th tree, orange nodes are in cut S

We perform the experiment of local pagerank, starting at the same node as **Nibble**. There is an extra reset probability a , and we set $a = 0.2$. The cut it produced cuts **989** edges and conductance **0.0323**. The shape of the cut is almost the same as in Figure 2. We also draw the first 4 levels of the 17th tree:

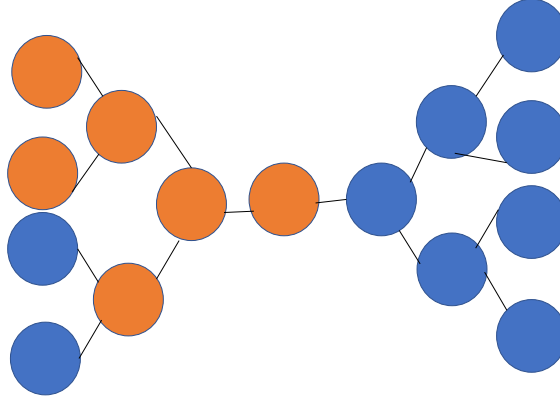


Figure 4: First 4 levels of 17^{th} tree, orange nodes are in cut S

	number of cut edges	conductance
Nibble	961	0.0314
local pagerank	989	0.0323

Table 2: Experiment summary of both algorithms

Notice both cuts produced by **Nibble** and local pagerank are worse than Cheeger’s cut, we have several conjectures about this result without formal proofs.

Conjecture 1 Cheeger’s cut uses only a **single** eigenvalue and eigenvector, so the cut structure is much simpler, on the other hand, both **Nibble** and local pagerank are variations of random walk, which means the vector we are working is a linear combination of multiple eigenvectors. Our previously thought was other eigenvectors will help to mitigate the problem caused by Cheeger’s cut, but through experiments, it seems the eigenstructure is not very easy to understand.

Conjecture 2 Notice the path size is a constant factor smaller than the desired size, which means due to Cheeger, the cut *should* look like a vertical cut, this might be a potential reason why the cut contains some small portion of lower trees.

4 Conclusions & Open Questions

We implemented modified versions of **Nibble** and local pagerank algorithms and experimented them on tree-cross-line graph, which is designed to be hard instance for these random walk based algorithms. Through experiments, we discovered that the cut these algorithms detected were not the best vertical cut nor the standard Cheeger-type cut, but a hybrid of both. We also have open questions for further research and work:

1. Our modified versions simplify the experiments, and by [1] and [2], it should only affect run time. However, it is possible that the cut quality is also affected by their remove of small entries, so one of the future directions is to implement vanilla version of these algorithms. Also, the path size in our experiment is a constant smaller than the desired size, so a more dedicated experiment on a “correct” size path will be performed.

2. Ultimately, tree-cross-line graph is a Cartesian product, so does there exist a class of algorithms that also has such product structure? For example, for random walk, can we simulate walks **independently** on tree and path, then combine the results of two walks in some way. We have been working simulating the most basic random walk on tree-cross-line graph in a product fashion, and hopefully can have an algorithm with a proof in this setting.
3. According to the experiment conducted by one of the coauthors Andy Yang, flow based algorithm by Wang et al. [9] does have a good performance on tree-cross-line graph, i.e., it can detect the best vertical cut on the graph. Can we work out a formal proof on why it works? Can we leverage more on the flow dynamics to develop more interesting algorithms on this problem?

Acknowledgment

We appreciate the help from Di Wang and Satish Rao on giving more clarifications and intuitions on their algorithm [9]. This work is a class project of the sequence of courses 15-300, Research and Innovation in Computer Science and 15-400, Research Practicum in Computer Science. We appreciate the advise of Bogdan Vasilescu and Jonathan Aldrich during the span of the project.

References

- [1] Daniel A. Spielman and Shang-Hua Teng. A Local Clustering Algorithm for Massive Graphs and its Application to Nearly-Linear Time Graph Partitioning. *arXiv e-prints*, page arXiv:0809.3232, September 2008.
- [2] Reid Andersen, Fan Chung, and Kevin Lang. Local graph partitioning using pagerank vectors. In *Proc IEEE Foundations of Computer Science*, pages 475–486, 10 2006.
- [3] Nathan Bronson, Zach Amsden, George Cabrera, Prasad Chakka, Peter Dimov, Hui Ding, Jack Ferris, Anthony Giardullo, Sachin Kulkarni, Harry Li, Mark Marchukov, Dmitri Petrov, Lovro Puzar, Yee Jiun Song, and Venkat Venkataramani. Tao: Facebook’s distributed data store for the social graph. pages 49–60, 01 2013.
- [4] William L. Hamilton, Rex Ying, and Jure Leskovec. Inductive representation learning on large graphs. In *Proceedings of the 31st International Conference on Neural Information Processing Systems*, NIPS’17, page 1025–1035, Red Hook, NY, USA, 2017. Curran Associates Inc.
- [5] Rex Ying, Ruining He, Kaifeng Chen, Pong Eksombatchai, William L. Hamilton, and Jure Leskovec. Graph convolutional neural networks for web-scale recommender systems. *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2018.
- [6] Reid Andersen and Yuval Peres. Finding sparse cuts locally using evolving sets. In *STOC ’09*, 2008.
- [7] László Lovász and Miklós Simonovits. Random walks in a convex body and an improved volume algorithm. *Random Struct. Algorithms*, 4:359–412, 1993.
- [8] Jeff Cheeger. A lower bound for the smallest eigenvalue of the laplacian. In *Proceedings of the Princeton conference in honor of Professor S. Bochner*, pages 195–199, 1969.

- [9] Di Wang, Kimon Fountoulakis, Monika Henzinger, Michael W. Mahoney, and Satish Rao. Capacity Releasing Diffusion for Speed and Locality. *arXiv e-prints*, page arXiv:1706.05826, June 2017.
- [10] A V Goldberg and R E Tarjan. A new approach to the maximum flow problem. In *Proceedings of the Eighteenth Annual ACM Symposium on Theory of Computing*, STOC '86, page 136–146, New York, NY, USA, 1986. Association for Computing Machinery.
- [11] Wei Qiao and Rifat Sipahi. Rules and limitations of building delay-tolerant topologies for coupled systems. *Physical review. E, Statistical, nonlinear, and soft matter physics*, 85:016104, 01 2012.

A Proof of Proposition 3.1

Proposition A.1. *The vector u can be computed directly via a Laplacian solve, specifically, u is the solution to the system*

$$(\beta D + L) u = \beta m$$

Where we assume $W = A^T D^{-1}$ to be the standard random walk transition matrix, and $\beta = \frac{a}{1-a}$.

Proof. Recall the system we originally wish to solve can be written as follows:

$$(I - (1 - a) W) p = am$$

Notice the vector u can be formed by a matrix-vector product: $u = D^{-1}p$, so substitute u and W in, we get

$$\begin{aligned} (I - (1 - a) A^T D^{-1}) D u &= am \\ (D - (1 - a) A) u &= am && A \text{ is symmetric} \\ (aD + (1 - a) D - (1 - a) A) u &= am \\ (aD + (1 - a) L) u &= am \\ (\beta D + L) u &= \beta m \end{aligned}$$

To see this reduces to a Laplacian solve, consider the following weighted graph G' : we add a ground node to G , call it v , and wire all nodes in G to v , with each edge weight be $\beta d(i)$ for edge $\{i, v\}$, and we can augment m by adding a -1 on its entry corresponding to v . This reduces the problem to solve a standard Laplacian system:

$$L' u' = m'$$

Then we can post-process vector u' by first subtracting multiples of $\mathbf{1}$, the all-one vector to make entry v being 0, then multiplying it by β . This gives our desired vector u . \square